

WEAREDEVELOPERS × BAND · DARK FACTORY · POCKETFUL TRACK · TEAM FOUR-SEATS

Four Seats Factory

Evidence Before Acceptance

Four autonomous agents built all four Pocketful stages from one human message — and accepted nothing until another agent had reproduced it.

Four seats, one rule

Each seat has a distinct, non-overlapping mandate. No seat accepts its own work. A claim never moves an artefact — a reproduction does.

Analyst

Turns the specification into numbered obligations before any code exists.

Implementer

The only seat that writes product code. All 21 product commits are its own.

Adversary

Attacks the plan, then the code, from its own clean clone and its own test suites.

Auditor

Reproduces everything in fresh containers. The only seat that can accept.

 **No seat accepts its own work. A claim never moves an artefact — a reproduction does.**

The mandates never name the product

Each seat runs on one generic standing instruction. Read it end to end: it never says payments, wallets, screens or this track.

"Build a numbered obligation register with precise source locations, expected behavior, boundaries, acceptance method and responsible seat. ... A sample suite is feedback, not the definition of the job."

— mandates/analyst-thgs.md

Point the factory at a different problem by changing the dispatch.

RUNBOOK.md says how.

Generic by design

The standing instructions contain no domain knowledge. The product domain lives only in the dispatch message.

Reusable factory

Change the dispatch, keep the four seats. The mandates carry no domain knowledge, so a new problem needs a new dispatch — not new mandates.

Obligations, not opinions

The analyst produces numbered, verifiable obligations — not prose summaries or vague goals.

One message, four stages

After the first message, nobody typed anything: the analyst opened every next stage itself.

UTC	Event
20:40	The only human message. It authorises all four stages.
21:35	Stage 1 accepted — no repair
22:46	Stage 2 accepted — after one rejection
00:08	Stage 3 accepted — after one rejection
00:43	Stage 4 accepted
00:55	All-stage verification from a fresh clone, independently checked

4h 14m

Total elapsed time

From the only human message to all-stage verification.

1

Human message

One dispatch authorised all four stages.
No further human input.

21

Product commits

All authored by the Implementer seat alone.

Every public check passed. It crashed anyway.

Stage-3 candidate 1, before it was rejected:

Check	Result
Official harness, isolated mode	147 + 35 + 6 — 0 skipped
Regression, API and UI at 375 and 1280 px	150 / 150
Migrations 1→3 and 2→3	14 / 14

Then: **heap out-of-memory after 4,761 statement reads. All state lost.**

The implementer had declared the risk and judged it acceptable. The analyst refused to take its word and sent both reviewers to probe it. **The auditor and the adversary crashed it independently.**

⊗ *A green check set is not evidence that a requirement holds.*

The repair, proved

	Rejected candidate	Accepted repair
Statement reads	crashed at 4,761	20,000 / 20,000 answered
Memory	1.06 GiB , then killed	39.7 MiB

Two candidates rejected

Three findings total — two by the adversary, one by the auditor. Every finding reproduced by a seat that did not write the code.

The smallest finding still blocked

A refused capture that silently reset the amount the user had typed. No money moved. It still blocked acceptance.


What it built

A wallet service: payments with idempotency, requests, bill splits, atomic settlements, holds with partial captures, historical balances, immutable corrections, refunds and correction batches. One Node.js process, **no dependencies**, every amount an exact integer.

Public checks per folder

147 · +35 · +6 · +5 — every stage claims its own

Live demo: four-seats-factory.onrender.com — `ada@example.com` / `correct horse`



Cost, and what we do not claim

~\$136 in model spend

Estimated at list prices — not a bill. More than 99% cache reads.

Verification costs more than building

The two verification seats cost more than the builder. That is where the money should go.

Public checks only

These are the public checks, not the hidden evaluation.

Three findings is not many

The auditor reproduces; it does not run mutation testing.

The run itself is the band's alone

The kit was prepared with ChatGPT and the documentation written with Claude Code. The run itself is the band's alone.

Every claim is mapped

github.com/hvaler/four-seats-factory

Every claim mapped to a file and a command in EVIDENCE-INDEX.md.

Methodology

- No seat accepts its own work
- Every finding reproduced independently
- A green check set is not evidence a requirement holds
- Reproduction, not assertion, is the standard of proof